

# Llama.cpp GPU Acceleration



Johannes Gäßler

# About Me

- Currently writing my master's thesis on experimental particle physics at KIT
- Additional bachelor's degree in informatics
- kafe2 (Karlsruhe Fit Environment 2) developer
- No prior experience with CUDA or language models

# llama.cpp

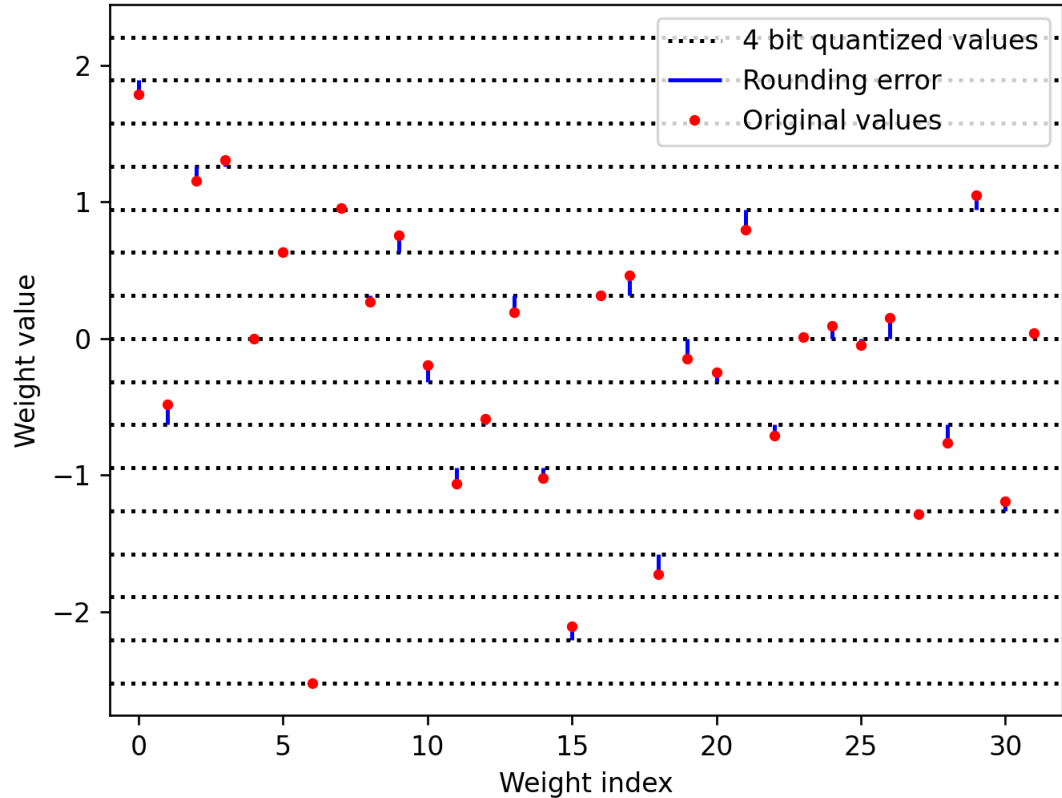
- Open-source C/C++ program for LLaMA inference



- Wide support across hardware and OSs
- Very good CPU performance
- I'm working on CUDA support


# Weight Quantization

- Original LLaMa weights are FP16
- Can be quantized to 4 bit ints with moderate quality loss
- int4 big model > FP16 small model



# Initial CUDA Implementation

- Dequantize weights to FP16/FP32, then use cuBLAS GEMM
- Performance only good for large batches
- Small batches: slower than CPU

[illegible]

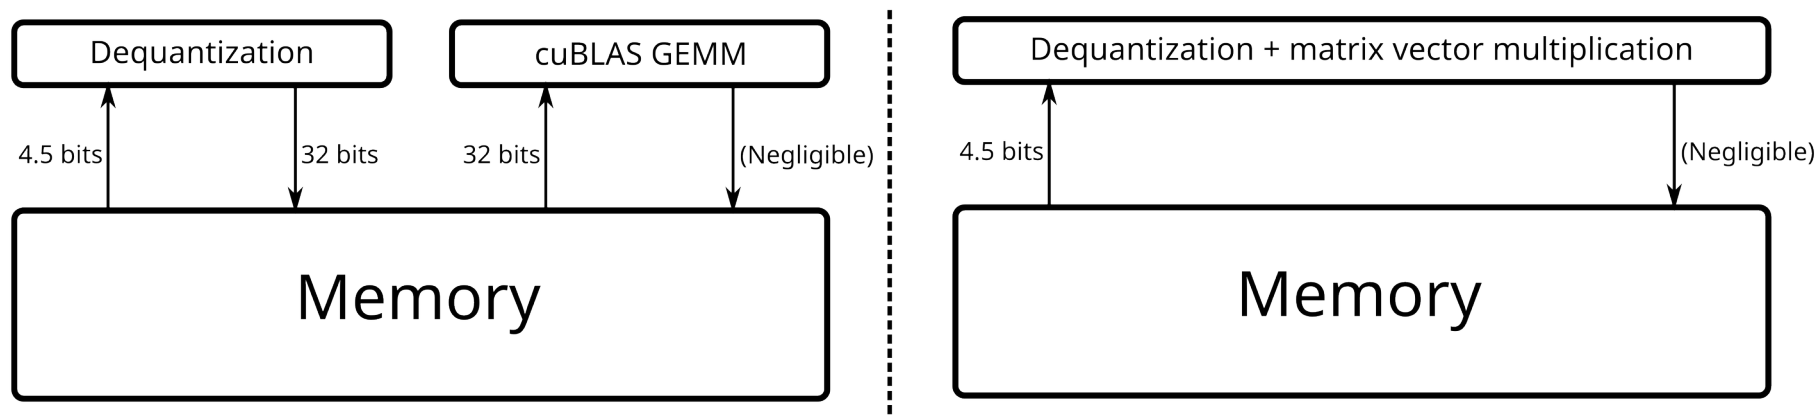
# Initial CUDA Implementation

- Matrix mult. I/O vs. compute depends on shape
- 2x square matrix:  $O(N^2)$  data,  $O(N^3)$  compute
- Square matrix + vector:  $O(N^2)$  data,  $O(N^2)$  compute
- Prompt processing: compute bound
- Token generation: I/O bound

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

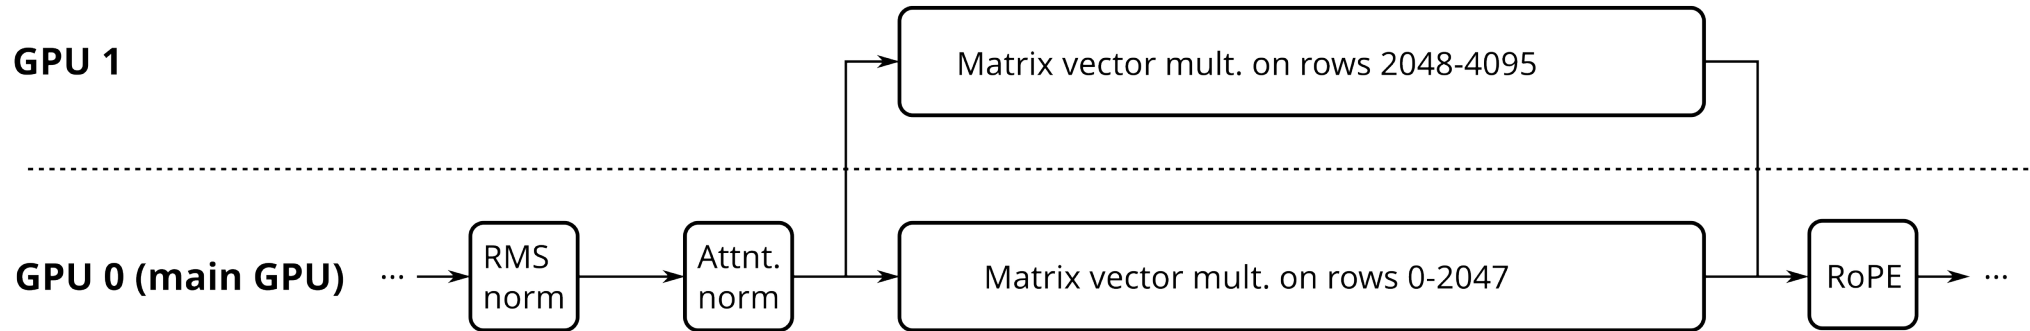
# Better CUDA Implementation

- Matrix vector multiplication + on-the-fly dequantization
- Per weight:  
 $36.5 \text{ bits read} + 32 \text{ bits write} \Rightarrow 4.5 \text{ bits read}$



# Better CUDA Implementation

- Multi GPU: split weight matrices across GPUs by rows
- Small tensors on main GPU only
- KV cache parallelization not implemented





# Current CUDA Implementation

- Don't dequantize matrix
- Instead quantize hidden state to 8 bit
- Use per-byte integer intrinsics (similar to CPU)

# Current CUDA Implementation

- $N$  blocks with 1 scale  $d$  and  $M$  values  $q_m$  each
- Dequantization:  $a_{inm} = d_{in}^a q_{inm}^a$ ,  $b_{nm} = d_n^b q_{nm}^b$

$$\begin{aligned} c_i &= \sum_{n=1}^N \sum_{m=1}^M a_{inm} b_{nm} = \sum_{n=1}^N \sum_{m=1}^M d_{in}^a q_{inm}^a d_n^b q_{nm}^b \\ &= \sum_{n=1}^N d_{in}^a d_n^b \sum_{m=1}^M q_{inm}^a q_{nm}^b. \end{aligned}$$

# Accessibility


- Open-source is a positive sum game
- Jevons Paradox: more efficiency => more use
- Main goal: reduce hardware requirements, get more users

# Accessibility: Memory

- Move weights from RAM to VRAM
- Can use total RAM + VRAM to fit model
- Goal: 33b q4 with standard settings on 16 GB RAM + 8 GB VRAM

# Accessibility: Quantization

- Hierarchical k-quants by I. Kawrakow
- 1 6 bit scale per 16 values
- 2 16 bit scales per 256 values
- Lower size/quantization error at the cost of more memory accesses



# Comparison to GPTQ

# LLaMA vs. LLaMA 2

# WIP: Quantized GEMM