

# Llama.cpp GPU Acceleration



Johannes Gäßler, 25.07.23

# About Me

- Currently writing my master's thesis on experimental particle physics at KIT
- Additional bachelor's degree in informatics
- kafe2 (Karlsruhe Fit Environment 2) developer
- No prior experience with CUDA or language models

# llama.cpp

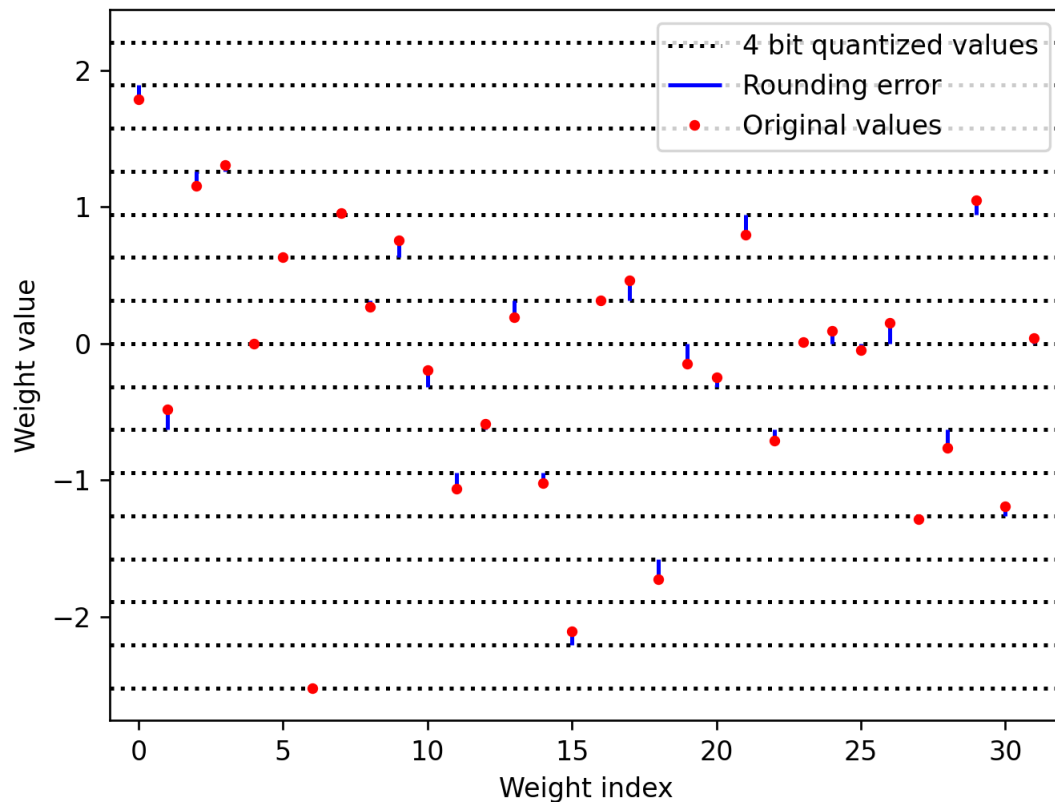
- Open-source C/C++ program for LLaMA inference



- Wide support across hardware and OSs
- Very good CPU performance
- I'm working on CUDA support


# Weight Quantization

- Original LLaMA weights are FP16
- Can be quantized to 4 bit ints with moderate quality loss
- int4 big model > FP16 small model



# Initial CUDA Implementation

- Dequantize weights to FP16/FP32, then use cuBLAS GEMM
- Performance only good for large batches
- Small batches: slower than CPU

[illegible]

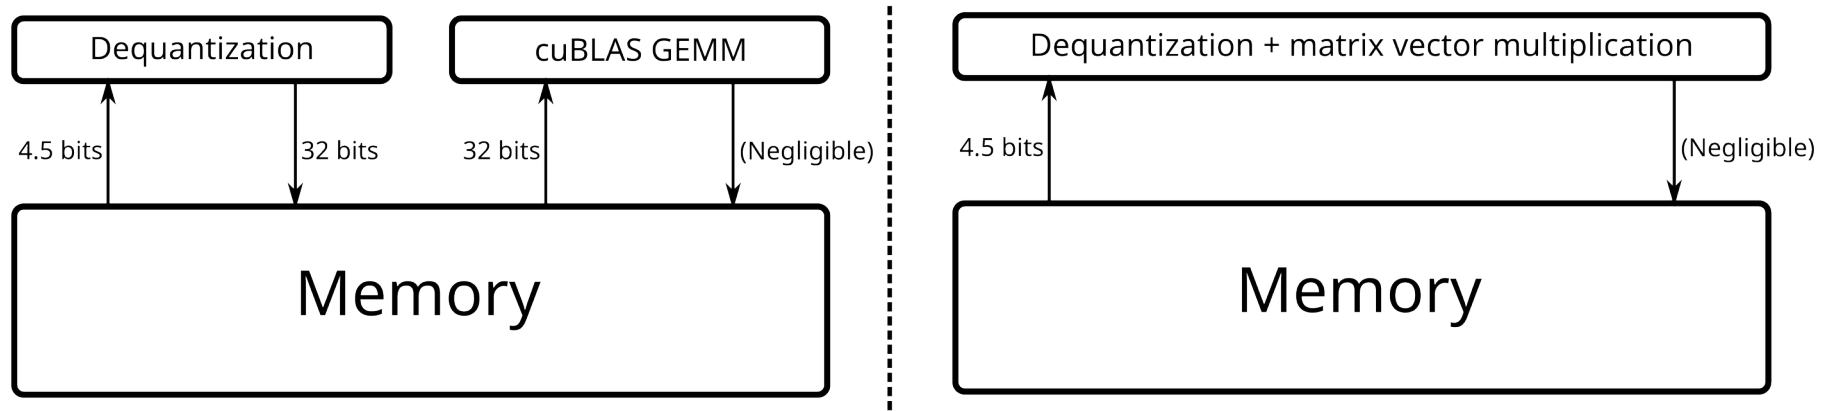
# Initial CUDA Implementation

- Matrix mult. I/O vs. compute depends on shape
- 2x square matrix:  $O(N^2)$  data,  $O(N^3)$  compute
- Square matrix + vector:  $O(N^2)$  data,  $O(N^2)$  compute
- Prompt processing: compute bound
- Token generation: I/O bound

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

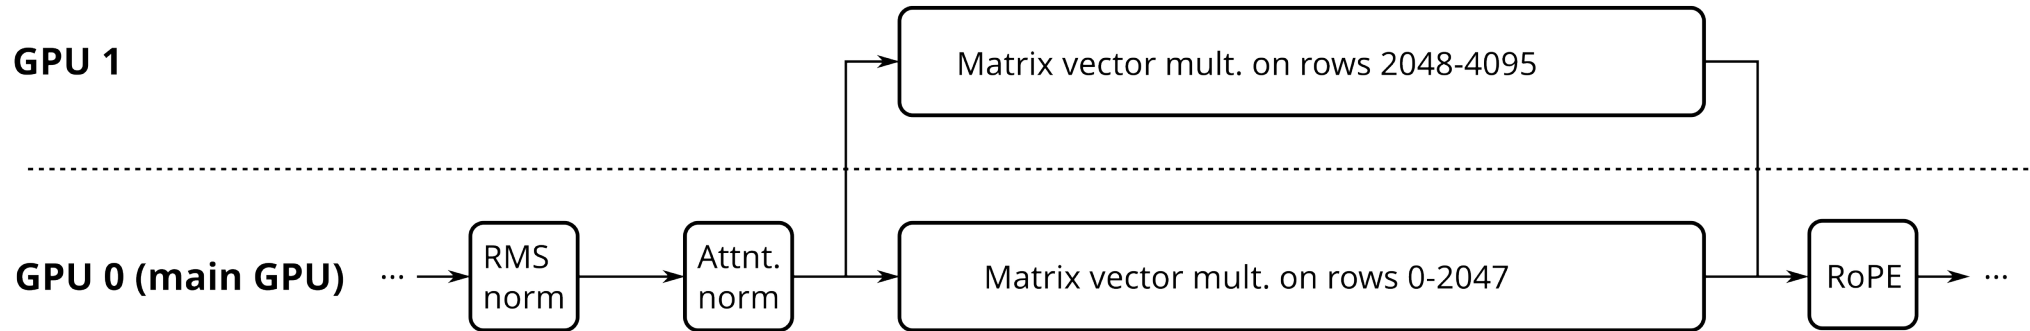
# Better CUDA Implementation

- Matrix vector multiplication + on-the-fly dequantization
- Per weight:  
 $36.5 \text{ bits read} + 32 \text{ bits write} \Rightarrow 4.5 \text{ bits read}$



# Better CUDA Implementation

- Multi GPU: split weight matrices across GPUs by rows
- Small tensors on main GPU only
- KV cache parallelization not implemented





# Current CUDA Implementation

- Don't dequantize matrix
- Instead quantize hidden state to 8 bit
- Use per-byte integer intrinsics (similar to CPU)

# Current CUDA Implementation

- $N$  blocks with 1 scale  $d$  and  $M$  values  $q_m$  each
- Dequantization:  $a_{inm} = d_{in}^a q_{inm}^a$ ,  $b_{nm} = d_n^b q_{nm}^b$

$$\begin{aligned} c_i &= \sum_{n=1}^N \sum_{m=1}^M a_{inm} b_{nm} = \sum_{n=1}^N \sum_{m=1}^M d_{in}^a q_{inm}^a d_n^b q_{nm}^b \\ &= \sum_{n=1}^N d_{in}^a d_n^b \sum_{m=1}^M q_{inm}^a q_{nm}^b. \end{aligned}$$

# Accessibility

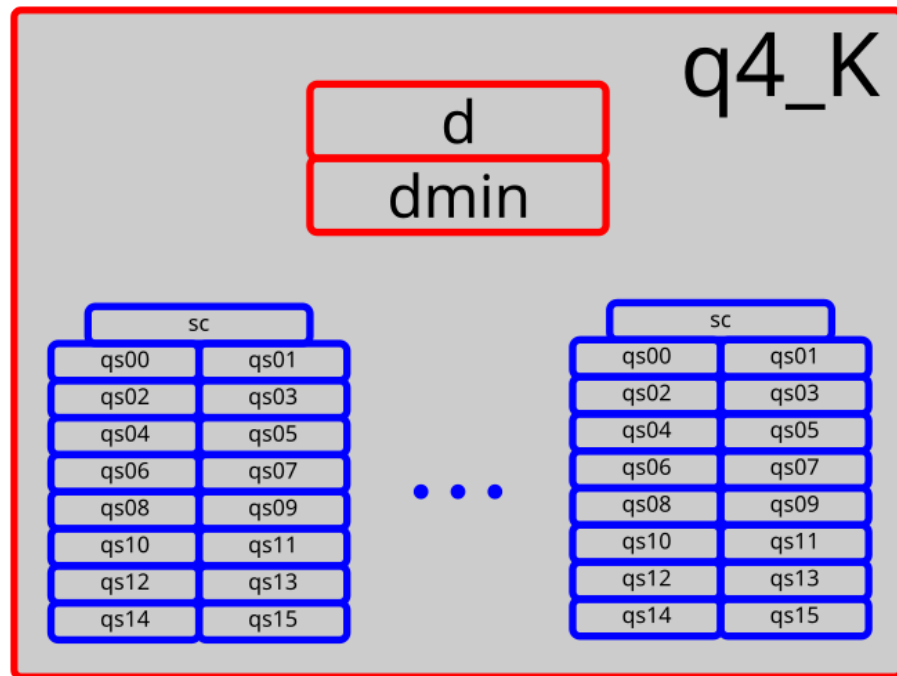
- Open-source is a positive sum game
- Jevons Paradox: more efficiency => more use
- Main goal: reduce hardware requirements, get more users

# Memory Optimization

- Move weights from RAM to VRAM
- Can use total RAM + VRAM to fit model
- Goal: 33b q4 with standard settings on 16 GB RAM + 8 GB VRAM

# k-quants by I. Kawrakow

- 1 6 bit scale per **block** (size 16)
- 2 16 bit scales per **superblock** (size 256)
- Lower size/quantization error at the cost of more memory accesses
- Different precision per tensor



# Comparison to GPTQ

Backend	Model	Prompt t/s	Gen. t/s	VRAM [MiB]	Perplexity**
AutoGPTQ (webui v1.3.1*)	7b q4 128g	Not measured	24.42	6028	6.54688
ExLlama (webui v1.3.1*)	7b q4 128g desc_act	<b>4076</b>	78.59	<b>5638</b>	6.28790
Llama.cpp (webui v1.3.1*)	7b q4_K_M	739	53.27	6960	<b>6.26391</b>
Llama.cpp (CLI rev. 84e09a7)	7b q4_K_M	965	<b>81.49</b>	6554	<b>6.26391</b>

Hardware: Ryzen 3700X, 32 GB RAM @ 3200 MHz, RTX 3090

\*<https://github.com/oobabooga/text-generation-webui/releases/tag/v1.3.1>

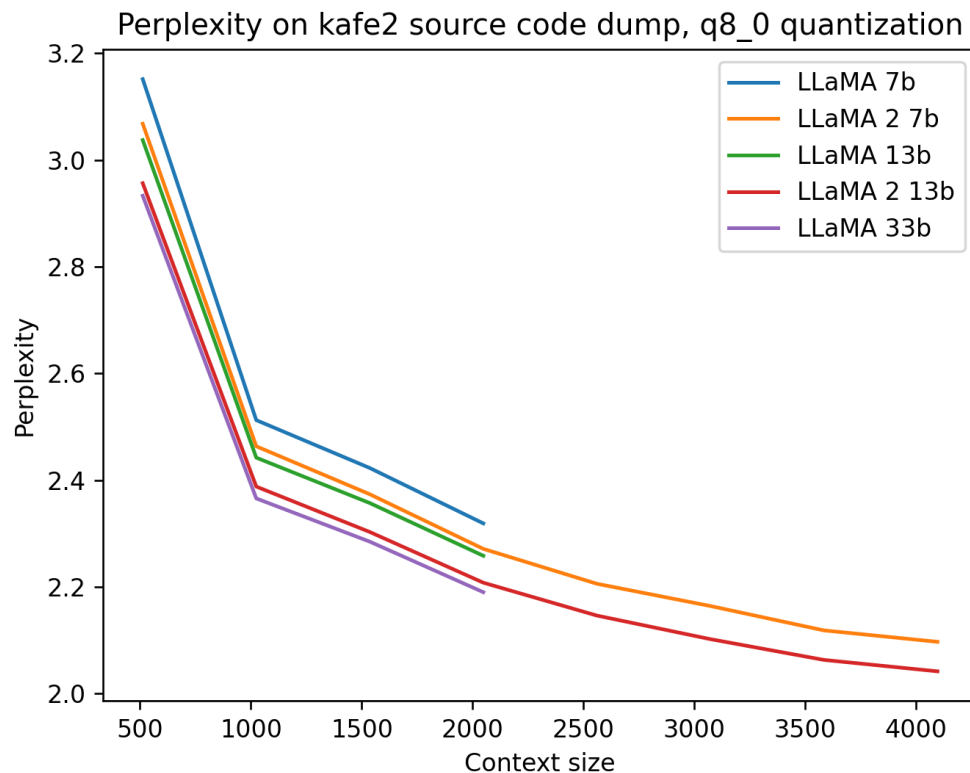
\*\*Perplexities taken from <https://oobabooga.github.io/blog/posts/perplexities/>

# WIP: Quantized GEMM

- Currently for batches: dequantize, then cuBLAS GEMM
- Try to write GEMM kernels that directly use quantized data
- Potentially faster and less VRAM usage
- Good GEMM kernels hard

# LLaMA vs. LLaMA 2


- 18.07.23: LLaMA 2 release
- Quasi open-source
- Trained on 2 trillion tokens (up from 1/1.4 trillion)
- Context: 2048 => 4096
- Grouped-query attention
- Chat finetunes included






Thank you for  
your attention!

# Appendix: Memory Scaling



# Appendix: Ryzen 3700X t/s

33b q4\_0, Ryzen 3700X, 3466 MHz dual channel RAM



# Appendix: Xeon E5-2683 v4 t/s

33b q4\_0, Xeon E5-2683 v4, 2133 MHz quad channel RAM

